



Publication-oriented architecture of the offline academic assistant pipeline

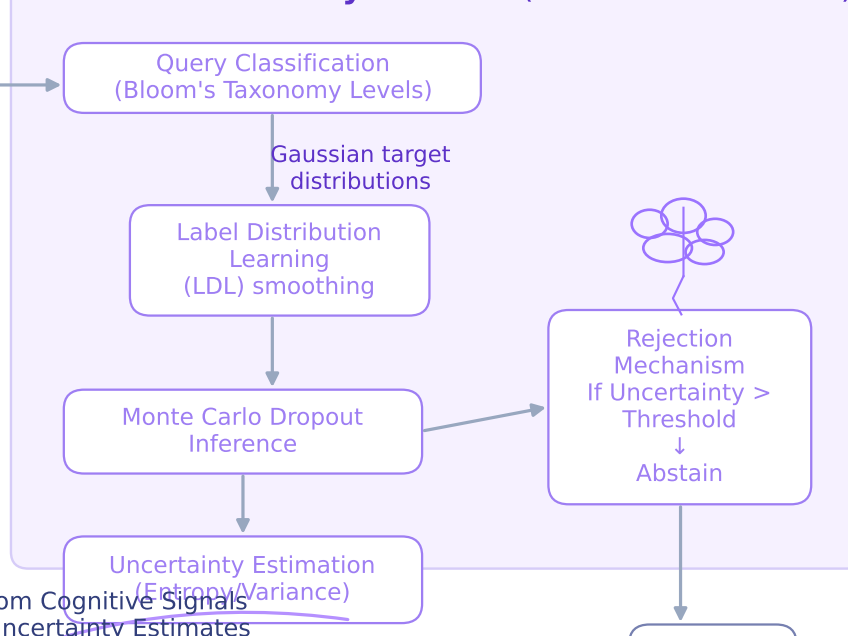
1. Multimodal Data Ingestion



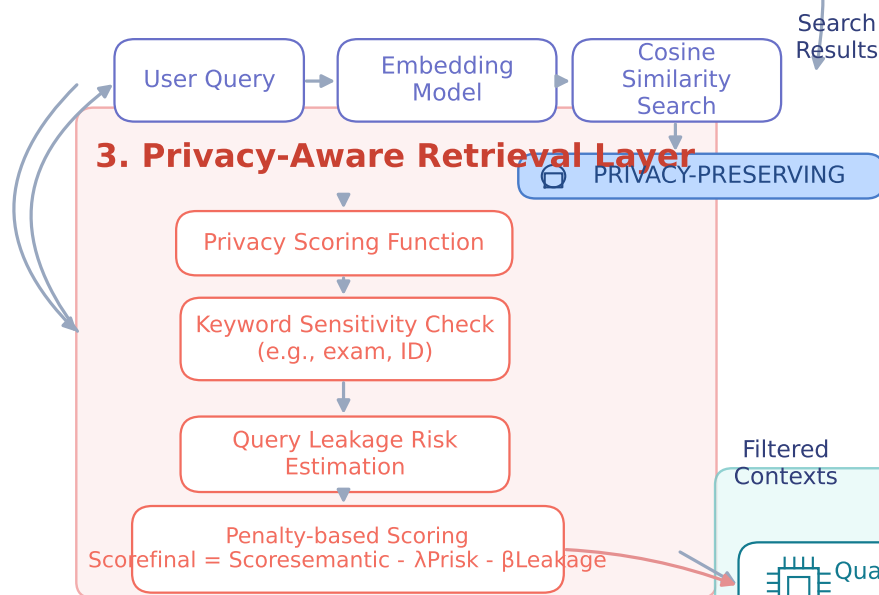
2. Vector-based RAG Engine



4. LDL Uncertainty Module (Monte Carlo based)



3. Privacy-Aware Retrieval Layer



5. Lightweight LLM-based generation

